

Objective: given i.i.d. samples, learn *to sample from* their distribution

Discrete Langevin Dynamics [1]

$$\mathbf{x}_t = \mathbf{x}_{t-1} + \frac{\epsilon}{2} \nabla_{\mathbf{x}} \log p(\mathbf{x}_{t-1}) + \sqrt{\epsilon} \mathbf{z}_t$$



Score Denoising Technique [3]

Explicit Score Matching $J^{\text{naive}}(\theta) = \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\|\mathbf{s}_{\theta}(\mathbf{x}) - \nabla_{\mathbf{x}} \log p_{\text{data}}(\mathbf{x})\|^2]$ Intractable because we don't know the log-density



Denoising Score Matching $J^{\text{denoising}}(\theta) = \frac{1}{2} \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x}), \tilde{\mathbf{x}} \sim q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})} [\|\mathbf{s}_{\theta}(\tilde{\mathbf{x}}) - \nabla_{\tilde{\mathbf{x}}} \log q_{\sigma}(\tilde{\mathbf{x}}|\mathbf{x})\|^2]$ SGD compatible

Continuous diffusion via SDEs [2]

$$d\mathbf{x}_t = \mathbf{f}(\mathbf{x}_t, t)dt + g(t)d\mathbf{w}_t$$

Controllable Generation [2]

Score-based framework is very flexible.

We just need to compute $\nabla_{\mathbf{x}} \log p_t(\mathbf{x}|\mathbf{y})$

We give two examples.

Reverse SDE

$$d\mathbf{x}_t = [\mathbf{f}(\mathbf{x}_t, t) - g(t)^2 \nabla_{\mathbf{x}} \log p_t(\mathbf{x}_t)]dt + d\tilde{\mathbf{w}}_t$$



Probability flow ODE

$$d\mathbf{x}_t = \left[\mathbf{f}(\mathbf{x}_t, t) - \frac{1}{2} g(t)^2 \nabla_{\mathbf{x}} \log p_t(\mathbf{x}_t) \right] dt$$



Class-conditional generation

$$\nabla_{\mathbf{x}} \log p_t(\mathbf{x}|\mathbf{y}) = \underbrace{\nabla_{\mathbf{x}} \log p_t(\mathbf{x})}_{\approx s_{\theta}(\mathbf{x}, t)} + \underbrace{\nabla_{\mathbf{x}} \log p_t(\mathbf{y}|\mathbf{x})}_{\text{estimated with backpropagation on time-conditional classifier } \hat{p}(\mathbf{y}|\mathbf{x}, t)}$$



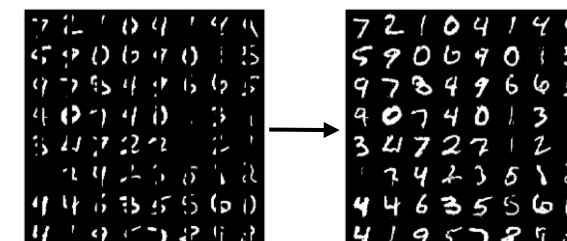
Imputation

$\mathbf{x}_0 \sim p_{\text{data}}(\mathbf{x})$
 $\Omega(\mathbf{x}_0)$ visible; $\bar{\Omega}(\mathbf{x}_0)$ masked

Reverse diffusion on the masked part

$$\nabla_{\bar{\Omega}(\mathbf{x}_t)} \log p_t([\bar{\Omega}(\mathbf{x}_t); \hat{\Omega}(\mathbf{x}_t)])$$

$$\hat{\Omega}(\mathbf{x}_t) \sim p_t(\Omega(\mathbf{x}_t) | \Omega(\mathbf{x}_0) = \mathbf{y})$$



[1] Yang Song and Stefano Ermon. *Generative Modeling by Estimating Gradients of the Data Distribution*. Advances in Neural Information Processing Systems 32 (NeurIPS), 2019.

[2] Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. *Score-Based Generative Modeling through Stochastic Differential Equations*. arXiv preprint arXiv:2011.13456, 2021.

[3] Pascal Vincent. *A Connection Between Score Matching and Denoising Autoencoders*. Neural Computation, 23(7):1661–1674, 2011.